

# Basic Machine Learning

## Supervised Learning

Kyunghyun Cho

New York University

Courant Institute (Computer Science) and Center for Data Science

Facebook AI Research

# Data-Driven Algorithm Design?

- Algorithm: a sequence of instructions that solves a problem.
- Traditionally,
  1. A detailed problem specification is given.
  2. An algorithm is designed to solve the described problem.
- Machine learning: data-driven algorithm design
  1. [A rough problem specification is given.] ← optional
  2. A set of examples is provided.\*
  3. A machine learning model is “trained” to solve the problem.
- A meta-approach to algorithm design

\* Examples could be provided in a variety of ways.<sup>2</sup>

# Supervised Learning – Overview

- Provided:
  1. a set of  $N$  input-output “training” examples
$$D = \{(x_1, y_1), \dots, (x_N, y_N)\}$$
  2. A per-example loss function
$$l(M(x), y) \geq 0$$
  3. Evaluation sets\*: validation and test examples  $D_{\text{val}}, D_{\text{test}}$
- What we must decide:
  1. Hypothesis sets  $\mathcal{H}_1, \dots, \mathcal{H}_M$ 
    - Each set consists of all compatible models
  2. Optimization algorithm

\* Often it is necessary to design a loss function.

\* Often these sets are created by holding out subsets of training examples.

# Supervised Learning – Overview

- Given:
  1.  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and  $D_{\text{val}}, D_{\text{test}}$
  2.  $l(M(x), y) \geq 0$
  3.  $\mathcal{H}_1, \dots, \mathcal{H}_M$
  4. Optimization algorithm
- Supervised learning finds an appropriate algorithm/model automatically
  1. For each hypothesis set  $\mathcal{H}_m$ , find the best model:\*

$$\hat{M}_m = \arg \min_{M \in \mathcal{H}_m} \sum_{n=1}^N l(M(x_n), y_n)$$

using the optimization algorithm.

# Supervised Learning – Overview

- Given:
  1.  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and  $D_{\text{val}}, D_{\text{test}}$
  2.  $l(M(x), y) \geq 0$
  3.  $\mathcal{H}_1, \dots, \mathcal{H}_M$
  4. Optimization algorithm
- Supervised learning finds an appropriate algorithm/model automatically
  1. [Training] For each hypothesis set  $\mathcal{H}_m$ , find the best model:\*

$$\hat{M}_m = \arg \min_{M \in \mathcal{H}_m} \sum_{n=1}^N l(M(x_n), y_n)$$

using the optimization algorithm and the **training set**.

# Supervised Learning – Overview

- Given:
  1.  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and  $D_{\text{val}}, D_{\text{test}}$
  2.  $l(M(x), y) \geq 0$
  3.  $\mathcal{H}_1, \dots, \mathcal{H}_M$
  4. Optimization algorithm
- Supervised learning finds an appropriate algorithm/model automatically
  2. [Model Selection]\* Among the trained models, select the best one

$$\hat{M} = \arg \min_{M \in \{\mathcal{H}_1, \dots, \mathcal{H}_M\}} \sum_{(x, y) \in D_{\text{val}}} l(M(x), y)$$

using the **validation set** loss.

\* If you're familiar with deep learning, “hyperparameter optimization” may be a more familiar term for you.

# Supervised Learning – Overview

- Given:
  1.  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and  $D_{\text{val}}, D_{\text{test}}$
  2.  $l(M(x), y) \geq 0$
  3.  $\mathcal{H}_1, \dots, \mathcal{H}_M$
  4. Optimization algorithm
- Supervised learning finds an appropriate algorithm/model automatically
  3. [Reporting] Report how well the best model *would* work

$$R(\hat{M}) \approx \frac{1}{|D_{\text{test}}|} \sum_{(x,y) \in D_{\text{test}}} l(\hat{M}(x), y)$$

using the **test set** loss.

# Supervised Learning – Overview

- Given:
  1.  $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$  and  $D_{\text{val}}, D_{\text{test}}$
  2.  $l(M(x), y) \geq 0$
  3.  $\mathcal{H}_1, \dots, \mathcal{H}_M$
  4. Optimization algorithm
- Supervised learning finds an appropriate algorithm/model automatically
- It results in an algorithm  $\hat{M}$  with an expected performance of  $R(\hat{M})$ .



# Supervised Learning


- Three points to consider both in research and in practice
  1. How do we decide/design a **hypothesis set**?
  2. How do we decide a **loss function**?
  3. How do we **optimize** the loss function?

# Hypothesis set – Neural Networks

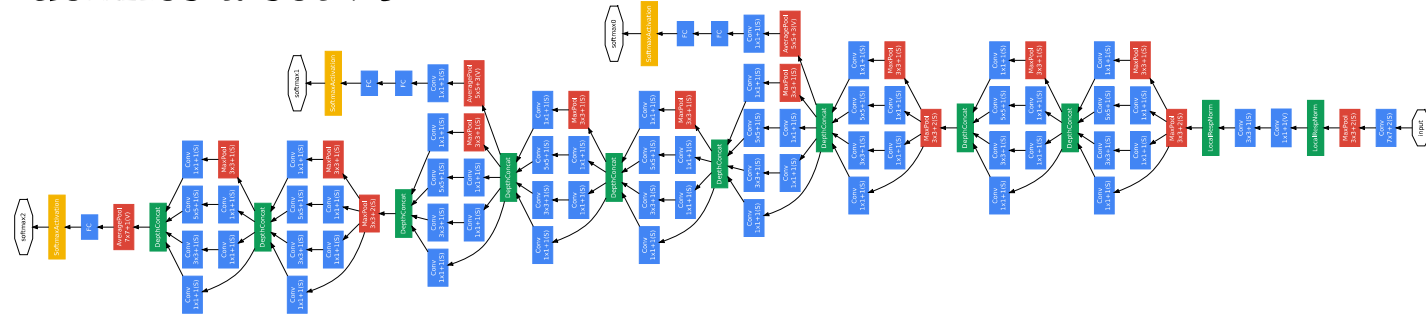
- What kind of machine learning approach will we consider?
  - Classification:
    - Support vector machines, Naïve Bayes classifier, logistic regression, ...?
  - Regression:
    - Support vector regression, Linear regression, Gaussian process, ...?
- How are the hyperparameters sets?
  - Support vector machines: regularization coefficient  $C$
  - Gaussian process: kernel function  $k(\cdot, \cdot)$

# Hypothesis set – Neural Networks

- In the case of deep learning/artificial neural networks,
  1. The architecture of a network defines a set  $\mathcal{H}$




vs.



2. Each model in the set  $M \in \mathcal{H}$  is characterized by its parameters  $\theta$ 
  - Weights and bias vectors define one model in the hypothesis set.
- There are infinitely many models in a hypothesis set.
- We use optimization to find “a” good model from the hypothesis set.

# Network Architectures

- What is a neural network? – An (arbitrary) directed acyclic graph (DAG)




1. Solid Circles  $\bigcirc$  : parameters (to be estimated or found)
2. Dashed Circles  $\bigcirc$  : vector inputs/outputs (given as a training example)
3. Squares  $\square$  : compute nodes (functions, often continuous/differentiable)

# Network Architectures


- What is a neural network? – An (arbitrary) directed acyclic graph (DAG)

1. Logistic regression

$$p_{\theta}(y = 1|x) = \sigma(w^{\top}x + b) = \frac{1}{1 + \exp(-w^{\top}x - b)}$$



2. 3<sup>rd</sup>-order polynomial function  $y = \alpha_0 + \alpha_1x + \alpha_2x^2 + \alpha_3x^3$




# Inference – Forward Computation

- What is a neural network? – An (arbitrary) directed acyclic graph (DAG)
- Forward computation: how you “use” a trained neural network.

- Topological sweep (breadth-first)
- Logistic regression

$$p_{\theta}(y = 1|x) = \sigma(w^{\top}x + b) = \frac{1}{1 + \exp(-w^{\top}x - b)}$$



# DAG $\leftrightarrow$ Hypothesis Set

- What is a neural network? – An (arbitrary) directed acyclic graph (DAG)
- Implication in practice
  - Naturally supports high-level abstraction
  - Object-oriented paradigm fits well.\*
    - Base classes: variable (input/output) node, operation node
    - Define the internal various types of variables and operations by inheritance
  - Maximal code reusability
    - See the success of PyTorch, TensorFlow, DyNet, Theano, ...
- You define a hypothesis set by designing a directed acyclic graph.
- The hypothesis space is then a set of all possible parameter settings.

# Supervised Learning

- Three points to consider both in research and in practice
  1. How do we decide/design a **hypothesis set**?
  2. How do we decide a **loss function**?
  3. How do we **optimize** the loss function?



# Loss Functions

- Per-example loss function
  - Computes how good a model is doing on a given example:
$$l(M(x), y) \geq 0$$
- So many loss functions...
  - Classification: hinge loss, log-loss, ...
  - Regression: mean squared error, mean absolute error, robust loss, ...
- In this lecture, we stick to distribution-based loss functions.

# Probability in 5 minutes – (1)

- An “event set”  $\Omega$  contains all possible events:  $\Omega = \{e_1, e_2, \dots, e_D\}$ 
  - Discrete: when there are a finite number of events  $|\Omega| < \infty$
  - Continuous: when there are infinitely many events  $|\Omega| = \infty$
- A “random variable”  $X$  could take any one of these events:  $X \in \Omega$
- A probability of an event:  $p(X = e_i)$ 
  - How likely would the  $i$ -th event happen?
  - How often has the  $i$ -th event occur relative to the other events?
- Properties
  1. Non-negative:  $p(X = e_i) \geq 0$
  2. Unit volume:  $\sum_{e \in \Omega} p(X = e) = 1$

# Probability in 5 minutes – (2)

- Multiple random variables: consider two here -  $X, Y$
- A joint probability  $p(Y = e_j^Y, X = e_i^X)$ 
  - How likely would  $e_j^Y$  and  $e_i^X$  happen together?
- A conditional probability  $p(Y = e_j^Y | X = e_i^X)$ 
  - Given  $e_i^X$ , how likely would  $e_j^Y$  happen?
  - The chance of both happening together divided by that of  $e_i^X$  happening regardless of whether  $e_j^Y$  happened:

$$p(Y|X) = \frac{p(X, Y)}{p(X)} \iff p(X, Y) = p(Y|X)p(X)$$

- Probability function  $p(X)$  returns a probability of  $X$

# Probability in 5 minutes – (3)

- Multiple random variables: consider two here -  $X, Y$
- A joint probability  $p(Y = e_j^Y, X = e_i^X)$ 
  - How likely would  $e_j^Y$  and  $e_i^X$  happen together?
- A marginal probability  $p(Y = e_j^Y)$ 
  - Regardless of what happens to  $X$ , how likely is  $e_j^Y$ ?

$$p(Y = e_j^Y) = \sum_{e \in \Omega_X} p(Y = e_j^Y, X = e)$$

# A Neural network

## computes a conditional distribution

- Supervised learning: what is  $y$  given  $x$ ?

$$f_{\theta}(x) = ?$$


- In other words, how probable is a certain value  $y'$  of  $y$  given  $x$ ?

$$p(y = y' | x) = ?$$


- What kind of distributions?
  - Binary classification: Bernoulli distribution
  - Multiclass classification: Categorical distribution
  - Linear regression: Gaussian distribution
  - Multimodal linear regression: Mixture of Gaussians

# Important distributions – Bernoulli

- How probable is a certain value  $y'$  of  $y$  given  $x$ ?  $p(y = y'|x) = ?$
- Binary classification: Bernoulli distribution  $\mathcal{B}(\mu)$ 
  - Probability:  $p(y|x) = \mu^y(1 - \mu)^{1-y}$ , where  $y \in \{0, 1\}$
  - Fully characterized by  $\mu \in [0, 1]$ .
  - A neural network then should turn the input  $x$  into  $\mu$




using a sigmoid function  $\sigma(a) = \frac{1}{1 + \exp(-a)}$



# Important distributions – Categorical

- How probable is a certain value  $y'$  of  $y$  given  $x$ ?  $p(y = y'|x) = ?$
- Multi-class classification: Categorical distribution  $\mathcal{C}(\{\mu_1, \mu_2, \dots, \mu_C\})$ 
  - Probability:  $p(y = v|x) = \mu_v$ , where  $\sum \mu_v = 1$
  - Fully characterized by  $\{\mu_1, \mu_2, \dots, \mu_C\}$ .
  - A neural network then should turn the input  $x$  into a vector  $\mu =$



$$\begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_C \end{bmatrix}$$

using a **softmax** function:  $\text{softmax}(a) = \frac{1}{\sum_{v=1}^C \exp(a_v)} \exp(a)$ .

# Important distributions – Gaussian

- How probable is a certain value  $y'$  of  $y$  given  $x$ ?  $p(y = y'|x) = ?$
- Regression: Gaussian distribution  $\mathcal{N}(\mu, \mathbb{I})$  with an identity covariance
  - Probability:  $p(y|x) = \frac{1}{Z} \exp(-\frac{1}{2}(y - \mu)^\top (y - \mu))$
  - Fully characterized by  $\mu \in \mathbb{R}^q$ .
  - A neural network then should turn the input  $x$  into a vector  $\mu$ .
  - Can be done trivially by affine transformation.



# Loss Function – negative log-probability

- Once a neural network outputs a conditional distribution  $p_\theta(y|x)$ , a natural way to define a loss function arises.
- Make sure training data is maximally likely:
  - Equiv. to making sure each and every training example is maximally likely.

$$\arg \max_{\theta} \log p_{\theta}(D) = \arg \max_{\theta} \sum_{n=1}^N \log p_{\theta}(y_n|x_n)$$

- Why *log*? – many reasons... but out of the lecture's scope.
- Equivalently, we want to minimize the *negative* log-probability.
  - A loss function is the sum of negative log-probabilities of correct answers.


$$L(\theta) = \sum_{n=1}^N l(M_{\theta}(x_n), y_n) = - \sum_{n=1}^N \log p_{\theta}(y_n|x_n)$$

# Loss Function – negative log-probability

- Once a neural network outputs a conditional distribution  $p_{\theta}(y|x)$ , a natural way to define a loss function arises.
- Practical implications
  - An OP node: negative log-probability (e.g., NLLLoss in PyTorch)
    - Inputs: the conditional distribution and the correct output
    - Output: the negative log-probability (a scalar)

# Loss Function – negative log-probability

- Once a neural network outputs a conditional distribution  $p_{\theta}(y|x)$ , a natural way to define a loss function arises.
- Logistic regression
  - Computes a Bernoulli distribution
  - Computes a negative log-probability
  - All in **one directed acyclic graph**
- Forward computation
  - Computes the conditional distribution, and
  - Computes the per-example loss



# Supervised Learning

- Three points to consider both in research and in practice
  1. How do we decide/design a **hypothesis set**?
  2. How do we decide a **loss function**?
  3. How do we **optimize** the loss function?

# Loss Minimization


- What we now know
  1. How to build a neural network with an arbitrary architecture.
  2. How to define a per-example loss as a negative log-probability.
  3. Define a single directed acyclic graph containing both.
- What we now need to know
  1. Choose an optimization algorithm.
  2. How to use the optimization algorithm to estimate parameters  $\theta$ .

# Local, iterative optimization

- An arbitrary function  $L : \mathbb{R}^d \rightarrow \mathbb{R}$
- Given the current value  $\theta_0$ , how should I move to minimize  $L$ ?
- Random guided search
  - Stochastically perturb  $\theta_0$ :  $\theta_k = \theta + \epsilon_k$ , where  $\epsilon_k \sim \mathcal{N}(0, s^2 \mathbf{1})$
  - Test each perturbed point  $L(\theta_k)$
  - Find the best perturbed point  $\theta_1 = \arg \min_{\theta_k} L(\theta_k)$
  - Repeat this until no improvement could be made.
- Applicable to any arbitrary loss function and neural network.
- Inefficient in the high-dimensional parameter space: large  $d$ .

# Local, iterative optimization

- An arbitrary function  $L : \mathbb{R}^d \rightarrow \mathbb{R}$
- Given the current value  $\theta_0$ , how should I move to minimize  $L$ ?
- Random guided search
  - Applicable to any arbitrary loss function and neural network.
  - Inefficient in the high-dimensional parameter space



# Gradient-based optimization


- A **continuous, differentiable**\* function  $L : \mathbb{R}^d \rightarrow \mathbb{R}$
- Given the current value  $\theta_0$ , how should I move to minimize  $L$ ?
- Gradient descent
  - The negative gradient of the function:  $-\nabla L(\theta_0)$
  - This is only valid in a local neighbourhood of  $\theta_0$ : take a very small step!
$$\theta = \theta_0 - \eta \nabla L(\theta_0)$$
- Efficient and effective even in the high dimensional space.
  - Can be improved with the second-order information (Hessian and/or FIM)

\* Almost everywhere, but not necessarily everywhere <sup>32</sup>



# Gradient-based optimization

- A **continuous, differentiable** function  $L : \mathbb{R}^d \rightarrow \mathbb{R}$
- Given the current value  $\theta_0$ , how should I move to minimize  $L$ ?
- Gradient descent
  - Efficient and effective even in the high dimensional space.
  - Learning rate must be carefully selected and annealed over time.



# Backward Computation – Backpropagation

- How do we compute the gradient of the loss function?

## 1. Manual derivation

- Relatively doable when the DAG is small and simple.
- When the DAG is larger and complicated, too much hassle.

## 2. Automatic differentiation (autograd)

- Use the chain rule of derivatives

$$\frac{\partial(f \circ g)}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x}$$

- The DAG is nothing but a composition of (mostly) differentiable functions.
- Automatically apply the chain rule of derivatives.

# Backward Computation – Backpropagation


- Automatic differentiation (autograd)
  1. Implement the Jacobian-vector product of each OP node:

$$\begin{bmatrix} \frac{\partial L}{\partial x_1} \\ \vdots \\ \frac{\partial L}{\partial x_d} \end{bmatrix} = \begin{bmatrix} \frac{\partial F_1}{\partial x_1} & \cdots & \frac{\partial F_{d'}}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial F_1}{\partial x_d} & \cdots & \frac{\partial F_{d'}}{\partial x_d} \end{bmatrix} \begin{bmatrix} \frac{\partial L}{\partial F_1} \\ \vdots \\ \frac{\partial L}{\partial F_{d'}} \end{bmatrix}$$

- Can be implemented efficiently without explicitly computing the Jacobian.
- The same implementation can be reused every time the OP node is called.


# Backward Computation – Backpropagation

- Automatic differentiation (autograd)
  2. Reverse-sweep the DAG starting from the loss function node.
    - Iteratively multiplies the Jacobian of each OP node until the leaf nodes of the parameters.
    - As expensive as forward computation with a constant overhead:  $O(N)$ , where  $N$ : # of nodes.



# Backward Computation – Backpropagation

- Practical Implications – Automatic differentiation (autograd)
  - Unless a complete new OP is introduced, no need to manually derive the gradient
  - Nice de-coupling of specification (front-end) and implementation (back-end)
    1. [Front-end] Design a neural network by creating a DAG.
    2. [Back-end] The DAG is “compiled” into an efficient code for a target compute device.



# Gradient-based Optimization

- Backpropagation gives us the gradient of the loss function w.r.t.  $\theta$
- Readily used by off-the-shelf gradient-based optimizers
  - Gradient descent, L-BFGS, Conjugate gradient, ...
  - Though, most are not applicable in a realistic neural network with 10s or 100s of millions of parameters.
- Stochastic gradient descent
  - Approximate the full loss function (the sum of per-examples losses) using only a small random subset of training examples:

$$\nabla L \approx \frac{1}{N'} \sum_{n=1}^{N'} \nabla l(M(x_{n'}), y_{n'})$$

# Stochastic Gradient Descent

- Stochastic gradient descent
  - Approximate the full loss function (the sum of per-examples losses) using only a small random subset of training examples:

$$\nabla L \approx \frac{1}{N'} \sum_{n=1}^{N'} \nabla l(M(x_{n'}), y_{n'})$$

- Unbiased estimate of the full gradient.\*
- Learning rate must be annealed appropriately.
- Extremely efficient *de facto* standard practice.

\* Under certain conditions

# Stochastic Gradient Descent

- Stochastic gradient descent in practice

1. Grab a random subset of  $M$  training examples\*

$$D' = \{(x_1, y_1), \dots, (x_{N'}, y_{N'})\}$$

2. Compute the minibatch gradient

$$\nabla L \approx \frac{1}{N'} \sum_{n=1}^{N'} \nabla l(M(x_{n'}), y_{n'})$$

3. Update the parameters

$$\theta \leftarrow \theta + \eta \nabla L(\theta; D')$$

4. Repeat until the validation loss stops improving.\*

\* In practice, sample without replacement until the training set is exhausted (one epoch).

\* This is called early-stopping which prevents the neural network from overfitting to training examples.




# Stochastic Gradient Descent – Early Stopping

- Stochastic gradient descent in practice
  1. Grab a random subset of  $M$  training examples
  2. Compute the minibatch gradient
  3. Update the parameters
  4. **Repeat until the validation loss stops improving.**
- An efficient way to prevent overfitting
  - Overfitting: the training loss is low, but the validation loss is not.
  - The most serious problem in statistical machine learning.
  - Early-stop based on the validation loss

# Stochastic Gradient Descent – Early Stopping

- An efficient way to prevent overfitting
  - Overfitting: the training loss is low, but the validation loss is not.
  - The most serious problem in statistical machine learning.
  - **Early-stop** based on the validation loss



# Stochastic Gradient Descent

## – Adaptive Learning Rate

- Stochastic gradient descent in practice
  1. Grab a random subset of  $M$  training examples  $D' = \{(x_1, y_1), \dots, (x_{N'}, y_{N'})\}$
  2. Compute the minibatch gradient
  3. Update the per-parameter learning rate  $\eta_\theta$
  4. Update the parameters
$$\theta \leftarrow \theta - \eta_\theta \frac{\partial L'}{\partial \theta}$$
  5. Repeat until the validation loss stops improving.
- Adaptive learning rate: Adam [Kingma&Ba, 2015], Adadelata [Zeiler, 2015], and many more...
  - Approximately re-scale parameters to improve the conditioning of the Hessian.

# Supervised Learning with Neural Networks

1. How do we decide/design a **hypothesis set**?
  - Design a network architecture as a directed acyclic graph
2. How do we decide a **loss function**?
  - Frame the problem as a conditional distribution modelling
  - The per-example loss function is a negative log-probability of a correct answer
3. How do we **optimize** the loss function?
  - Automatic backpropagation: no manual gradient derivation
  - Stochastic gradient descent with early stopping [and adaptive learning rate]

In the next lecture,

- We will study the entire cycle for text classification with a neural network.